

```

pip install pymongo

Requirement already satisfied: pymongo in /usr/local/lib/python3.10/dist-packages (4.4.1)
Requirement already satisfied: dnspython<3.0.0,>=1.16.0 in /usr/local/lib/python3.10/dist-packages (from pymongo) (2.4.2)

import pandas as pd
import numpy as np
import pymongo
import json

from google.colab import drive
drive.mount('/content/drive/')

!ls '/content/drive/My Drive/'
filepath = '/content/drive/My Drive/'

Drive already mounted at /content/drive/; to attempt to forcibly remount, call drive.mount("/content/drive/", force_remount=True).
'01a. UrbanSEC.gdoc'
'01b. RURAL SEC GRID.gdoc'
'01c. New SEC Qns and Grid.gdoc'
'01. SAMPLING DESIGN.gslides'
'02. Solution - Sample Size determination.gsheet'
'03a. Seat Computation (1).gsheet'
'03a. Seat Computation.gsheet'
'03. Opinion Polls and its complexities.gslides'
4.gsheet
'AbhishekLal_C23039_Data Exploration in Spark.ipynb'
'Abhishek Lal_C23039.xlsx'
'AbhishekLal_Resume (1).pdf'
AbhishekLal_Resume.pdf
'Adventure Works Sales (Multiple Sheets).gsheet'
advertising.csv
'AIF 01 nifty cos data.gsheet'
'AIF 01 two asset portfolio (1).gsheet'
'AIF 01 two asset portfolio.gsheet'
'AIF 02 BANK DATA.gsheet'
bangladesh.gsheet
'car data.csv'
cardekho.gsheet
'Case - Data Driven Decision Making.gslides'
Classroom
'Class Work Dataset(Abhishek Lal) (1).gsheet'
'Class Work Dataset(Abhishek Lal).gsheet'
'class work.gsheet'
'Colab Notebooks'
communication_infrastructure.gsheet
'data for class 17th.gsheet'
'Data variable.txt'
'FSA DuPont-Apollo tyres.gsheet'
GapminderHealth.gsheet
GTAS_review.txt
gta_data_batch_2.csv
kddcup.data_10_percent_corrected
mean_years_school.gsheet
"Participants' Worksheet.gsheet"
PCA.ipynb
'PGPDS Student Abhishek Lal.jpg'
pima-indians-diabetes.gsheet
PX03_Churn.gsheet
'PX22DS-Introduction_to_ML-05_Classification_KNN (1).gdoc'
PX22DS-Introduction_to_ML-05_Classification_KNN.gdoc
Python_Programs_Set_1
Python_Programs_Set_2
san_francisco.gsheet
'Supermarket aggregate data.csv'
Train_BigMart.gsheet
'Uber Cab Fare.csv'
UNHealth.gsheet
weatherAUS.csv
weatherAUS.gsheet

xy = pd.read_csv(filepath + 'Uber Cab Fare.csv')

```

▸ Reading the Data and Converting the Same to Dictionary

```

data=xy.to_dict(orient="records")
data

```

```

    'num_of_passengers': 1.0,
    'fare': 116.25,
    'tip': 0,
    'miscellaneous_fees': 6.0,
    'total_fare': 122.25,
    'surge_applied': 0},
    {'trip_duration': 955.0,
    'distance_traveled': 17.24,
    'num_of_passengers': 1.0,
    'fare': 221.25,
    'tip': 57,
    'miscellaneous_fees': 5.850000000000023,
    'total_fare': 284.1,
    'surge_applied': 0},
    {'trip_duration': 2687.0,
    'distance_traveled': 23.64,
    'num_of_passengers': 1.0,
    'fare': 345.0,
    'tip': 94,
    'miscellaneous_fees': 30.200000000000102,
    'total_fare': 469.2000000000001,
    'surge_applied': 1},
    {'trip_duration': 415.0,
    'distance_traveled': 1.87,
    'num_of_passengers': 1.0,
    'fare': 48.75,
    'tip': 0,
    'miscellaneous_fees': 26.625,
    'total_fare': 75.375,
    'surge_applied': 1},
    {'trip_duration': 654.0,
    'distance_traveled': 2.59,
    'num_of_passengers': 1.0,
    'fare': 67.5,
    'tip': 15,
    'miscellaneous_fees': 5.700000000000003,
    'total_fare': 88.2,
    'surge_applied': 0},
    {'trip_duration': 2251.0,
    'distance_traveled': 10.65,
    'num_of_passengers': 1.0,
    'fare': 198.75,
    'tip': 42,
    'miscellaneous_fees': 13.949999999999989,
    'total_fare': 254.7,
    'surge_applied': 0},
    {'trip_duration': 2540.0,
    'distance_traveled': 12.23,
    'num_of_passengers': 1.0,
    'fare': 243.75,
    'tip': 0,
    'miscellaneous_fees': 13.5,
    'total_fare': 257.25,
    'surge_applied': 0},
    ...]

```

▼ Installation of Spark

```

!pip3 -q install pyspark
from pyspark.sql import SparkSession
spark = SparkSession.builder.appName('Praxis').getOrCreate()

from pyspark.sql import SparkSession
from pyspark.conf import SparkConf
from pyspark.sql.types import *
import pyspark.sql.functions as F
from pyspark.sql.functions import col, asc, desc
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from pyspark.sql import SQLContext
from pyspark.mllib.stat import Statistics
import pandas as pd
from pyspark.sql.functions import udf
from pyspark.ml.feature import OneHotEncoder, StringIndexer, VectorAssembler, StandardScaler
from pyspark.ml import Pipeline
from sklearn.metrics import confusion_matrix

```

▼ Converting MongoDB data to spark dataframe

```
ls
```

```
drive/ sample_data/
```

SCHEMA FOR SPARK TABLE

```
# schema='{ "fields": [{ "metadata": {}, "name": "selling_price", "nullable": true, "type": "long" }, { "metadata": {}, "name": "km_driven", "nullable": tr

# new_schema = StructType.fromJson(json.loads(schema))
```

▼ Load Data

```
# Load and Read the dataset
data=spark.read.csv('/content/drive/My Drive/Uber Cab Fare.csv',inferSchema=True,header=True)

# Chech the datatypes of the inputs
data.printSchema()

root
|-- trip_duration: double (nullable = true)
|-- distance_traveled: double (nullable = true)
|-- num_of_passengers: double (nullable = true)
|-- fare: double (nullable = true)
|-- tip: integer (nullable = true)
|-- miscellaneous_fees: double (nullable = true)
|-- total_fare: double (nullable = true)
|-- surge_applied: integer (nullable = true)
```

▼ Data Exploration

```
data.count(), len(data.columns)
```

```
(209673, 8)
```

```
data.dtypes
```

```
[('trip_duration', 'double'),
 ('distance_traveled', 'double'),
 ('num_of_passengers', 'double'),
 ('fare', 'double'),
 ('tip', 'int'),
 ('miscellaneous_fees', 'double'),
 ('total_fare', 'double'),
 ('surge_applied', 'int')]
```

```
data.printSchema()
```

```
root
|-- trip_duration: double (nullable = true)
|-- distance_traveled: double (nullable = true)
|-- num_of_passengers: double (nullable = true)
|-- fare: double (nullable = true)
|-- tip: integer (nullable = true)
|-- miscellaneous_fees: double (nullable = true)
|-- total_fare: double (nullable = true)
|-- surge_applied: integer (nullable = true)
```

```
data.count()
```

```
209673
```

```
data.columns
```

```
['trip_duration',
 'distance_traveled',
 'num_of_passengers',
 'fare',
 'tip',
 'miscellaneous_fees',
 'total_fare',
 'surge_applied']
```

```
data.describe().show()
```

summary	trip_duration	distance_traveled	num_of_passengers	fare	tip	miscellaneous_fees	to
count	209673	209673	209673	209673	209673	209673	209673
mean	1173.181477824994	5.054431185703426	1.2928083253447034	99.62343112847154	13.030824188140581	15.13682865700149	127.791083
stddev	4775.653620869363	125.21741931669423	0.930753593324805	85.60270241033903	20.36776412100705	12.553435595551546	98.797374
min	0.0	0.02	0.0	0.0	0	-0.5	
max	86387.0	57283.91	9.0	4466.25	2500	435.0	

```
for col in data.columns:
    print(col, "has", data.filter(data[col].isNull()).count(), "Null values.")
```

trip_duration has 0 Null values.
distance_traveled has 0 Null values.
num_of_passengers has 0 Null values.
fare has 0 Null values.
tip has 0 Null values.
miscellaneous_fees has 0 Null values.
total_fare has 0 Null values.
surge_applied has 0 Null values.

```
data.where(data.surge_applied == 0).show()
```

trip_duration	distance_traveled	num_of_passengers	fare	tip	miscellaneous_fees	total_fare	surge_applied
748.0	2.75	1.0	75.0	24	6.299999999999997	105.3	0
1187.0	3.43	1.0	105.0	24	13.200000000000001	142.20000000000002	0
671.0	5.63	3.0	90.0	0	9.75	99.75	0
329.0	2.09	1.0	45.0	12	13.199999999999999	70.19999999999999	0
453.0	2.22	1.0	52.5	0	6.0	58.5	0
134.0	1.48	1.0	33.75	0	6.0	39.75	0
980.0	3.48	1.0	90.0	0	6.0	96.0	0
305.0	2.04	1.0	45.0	0	9.75	54.75	0
4059.0	30.0	1.0	390.0	0	55.125	445.125	0
530.0	2.43	1.0	56.25	0	9.75	66.0	0
773.0	3.41	2.0	75.0	9	13.5	97.5	0
357.0	1.43	1.0	112.5	0	2.25	114.75	0
1163.0	5.62	1.0	108.75	0	13.5	122.25	0
1059.0	5.39	1.0	112.5	0	6.0	118.5	0
1467.0	5.21	1.0	127.5	0	6.0	133.5	0
929.0	4.31	1.0	93.75	25	5.9000000000000006	124.65	0
1023.0	2.41	2.0	90.0	19	6.125	115.125	0
255.0	2.48	1.0	45.0	0	6.0	51.0	0
668.0	3.11	2.0	75.0	0	9.75	84.75	0
11323.0	0.14	1.0	337.5	68	2.1999999999999886	407.7	0

only showing top 20 rows

```
new_data = data.filter((data.surge_applied != 0))
```

```
new_data.show()
```

trip_duration	distance_traveled	num_of_passengers	fare	tip	miscellaneous_fees	total_fare	surge_applied
730.0	3.12	1.0	71.25	0	26.625	97.875	1
355.0	1.74	1.0	45.0	0	26.625	71.625	1
1288.0	5.21	1.0	116.25	36	26.324999999999999	178.575	1
84305.0	4.94	2.0	116.25	29	26.199999999999999	171.45	1
328.0	2.16	1.0	48.75	15	26.700000000000003	90.45	1
1848.0	4.83	1.0	142.5	17	26.425000000000001	185.925	1
234.0	1.61	0.0	37.5	0	26.625	64.125	1
1537.0	7.64	1.0	146.25	14	26.425000000000001	186.675	1
319.0	1.67	1.0	41.25	0	26.625	67.875	1
375.0	1.45	1.0	45.0	24	33.825	102.825	1
586.0	2.57	1.0	63.75	0	26.625	90.375	1
1617.0	3.23	6.0	123.75	8	26.125	157.875	1
1271.0	4.84	1.0	116.25	0	26.625	142.875	1
915.0	4.46	1.0	93.75	26	33.700000000000002	153.45000000000002	1
742.0	4.89	3.0	90.0	11	30.625	131.625	1
1149.0	9.35	1.0	153.75	0	26.625	180.375	1
1146.0	3.88	1.0	97.5	37	26.825000000000001	161.32500000000002	1
3722.0	19.17	1.0	322.5	89	34.300000000000001	445.8	1
1911.0	6.76	1.0	161.25	47	26.574999999999999	234.825	1
622.0	2.03	1.0	60.0	0	26.625	86.625	1

only showing top 20 rows

```
from pyspark.sql.functions import *
```

Showing group wise average selling_price

```
data.groupBy('surge_applied').mean('total_fare').show()
```

+-----+-----+	
surge_applied	avg(total_fare)
+-----+-----+	
1	170.0919412425104
0	111.29808193296972
+-----+-----+	

```
import pandas as pd
import numpy as np
```

```
DF = data.toPandas()
```

```
DF.head()
```

	trip_duration	distance_traveled	num_of_passengers	fare	tip	miscellaneous_fees	total_fare	surge_applied
0	748.0	2.75	1.0	75.00	24	6.300	105.300	0
1	1187.0	3.43	1.0	105.00	24	13.200	142.200	0
2	730.0	3.12	1.0	71.25	0	26.625	97.875	1
3	671.0	5.63	3.0	90.00	0	9.750	99.750	0
4	329.0	2.09	1.0	45.00	12	13.200	70.200	0

```
cor = DF.corr()
```

```
cor.total_fare.sort_values(ascending=False)
```

total_fare	1.000000
fare	0.966748
tip	0.508639
miscellaneous_fees	0.452568
surge_applied	0.267350
trip_duration	0.142159
distance_traveled	0.036677
num_of_passengers	0.014234
Name: total_fare, dtype: float64	

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
sns.distplot(x=DF['total_fare'], bins=20, kde=True)
```

```
<ipython-input-44-b42c81b330c0>:1: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.boxplot(x='total_fare', data=DF)
```

```
<Axes: xlabel='total_fare'>
```



```
sns.boxplot(x='distance_traveled', data=DF)
```

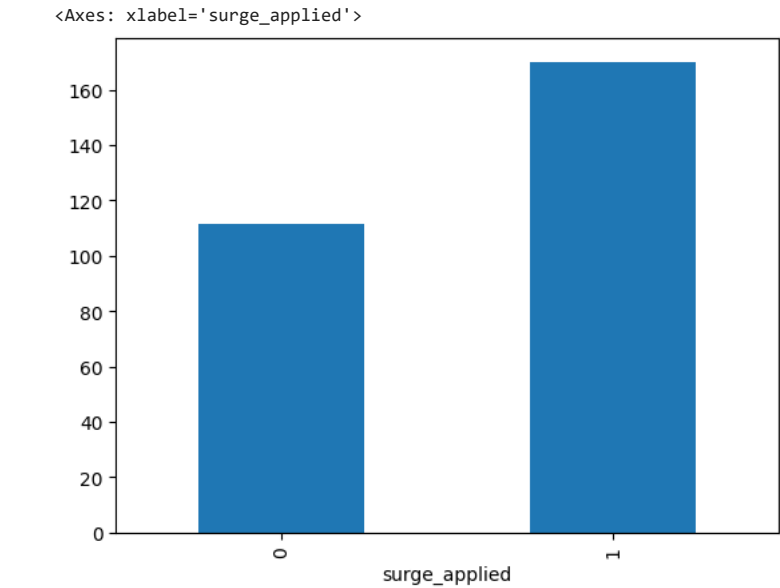
```
<Axes: xlabel='distance_traveled'>
```




```
DF['total_fare'].plot(kind='hist', bins=25)
```



```
plt.figure()
DF.groupby('surge_applied')['total_fare'].mean().plot(kind='bar')
```



```
sns.scatterplot(x='total_fare', y='fare',data=DF )
```



```
from datetime import date
```

```
data.show(5)
```

trip_duration	distance_traveled	num_of_passengers	fare	tip	miscellaneous_fees	total_fare	surge_applied
748.0	2.75	1.0	75.0	24	6.2999999999999997	105.3	0
1187.0	3.43	1.0	105.0	24	13.2000000000000017	142.20000000000002	0
730.0	3.12	1.0	71.25	0	26.625	97.875	1
671.0	5.63	3.0	90.0	0	9.75	99.75	0
329.0	2.09	1.0	45.0	12	13.199999999999989	70.19999999999999	0

only showing top 5 rows

▼ Checking For Correlation

```
from pyspark.sql.functions import corr

type(data)

pyspark.sql.dataframe.DataFrame
```

▼ Data Preparation using One-Hot Encoder and Vector Assembler

```
from pyspark.ml.feature import StringIndexer

from pyspark.ml.feature import OneHotEncoder

from pyspark.ml.feature import VectorAssembler

data.show(7)
```

trip_duration	distance_traveled	num_of_passengers	fare	tip	miscellaneous_fees	total_fare	surge_applied
748.0	2.75	1.0	75.0	24	6.299999999999997	105.3	0
1187.0	3.43	1.0	105.0	24	13.200000000000001	142.20000000000002	0
730.0	3.12	1.0	71.25	0	26.625	97.875	1
671.0	5.63	3.0	90.0	0	9.75	99.75	0
329.0	2.09	1.0	45.0	12	13.199999999999989	70.19999999999999	0
355.0	1.74	1.0	45.0	0	26.625	71.625	1
453.0	2.22	1.0	52.5	0	6.0	58.5	0

only showing top 7 rows

```
from pyspark.ml.linalg import DenseVector

from pyspark.ml.feature import StandardScaler

input_data = new_data.rdd.map(lambda x: (x[0], DenseVector(x[1:])))

df_assembler = VectorAssembler(inputCols=['trip_duration', 'distance_traveled', 'num_of_passengers', 'tip', 'surge_applied'], outputCol="features")

#df = df_assembler.transform(df)

df = df_assembler.transform(data)

df.show(5)
```

trip_duration	distance_traveled	num_of_passengers	fare	tip	miscellaneous_fees	total_fare	surge_applied	features
748.0	2.75	1.0	75.0	24	6.299999999999997	105.3	0	[748.0,2.75,1.0,2...
1187.0	3.43	1.0	105.0	24	13.200000000000001	142.20000000000002	0	[1187.0,3.43,1.0,..
730.0	3.12	1.0	71.25	0	26.625	97.875	1	[730.0,3.12,1.0,0...
671.0	5.63	3.0	90.0	0	9.75	99.75	0	[671.0,5.63,3.0,0...
329.0	2.09	1.0	45.0	12	13.199999999999989	70.19999999999999	0	[329.0,2.09,1.0,1...

only showing top 5 rows

```
df.select(['features','total_fare']).show(5)
```

features	total_fare
[748.0,2.75,1.0,2...	105.3
[1187.0,3.43,1.0,..	142.20000000000002
[730.0,3.12,1.0,0...	97.875
[671.0,5.63,3.0,0...	99.75
[329.0,2.09,1.0,1...	70.19999999999999

only showing top 5 rows

▼ Building and Comparing ML Models

▼ Creation and application of 2 Transformers / Estimators in a pipeline

```
# pipeline_stages=Pipeline()\n#\n# pipeline_model=pipeline_stages.fit(new_data)\n\nfrom pyspark.ml.feature import StandardScaler\n\n# Initialize the `standardScaler`\nstandardScaler = StandardScaler(inputCol="features", outputCol="features_scaled")\n\n# Fit the DataFrame to the scaler\nscaler = standardScaler.fit(df)\n\n# Transform the data in `df` with the scaler\nscaled_df = scaler.transform(df)\n\nscaled_df.take(2)\n\n[Row(trip_duration=748.0, distance_traveled=2.75, num_of_passengers=1.0, fare=75.0, tip=24, miscellaneous_fees=6.299999999999997,\ntotal_fare=105.3, surge_applied=0, features=DenseVector([748.0, 2.75, 1.0, 24.0, 0.0])), features_scaled=DenseVector([0.1566,\n0.022, 1.0744, 1.1783, 0.0])),\n Row(trip_duration=1187.0, distance_traveled=3.43, num_of_passengers=1.0, fare=105.0, tip=24,\nmiscellaneous_fees=13.2000000000000017, total_fare=142.20000000000002, surge_applied=0, features=DenseVector([1187.0, 3.43, 1.0,\n24.0, 0.0])), features_scaled=DenseVector([0.2486, 0.0274, 1.0744, 1.1783, 0.0]))]
```

Train-Test Split

```
# Split the data into train and test sets\ntrain_data, test_data = scaled_df.randomSplit([.7,.3],seed=1234)\nprint("Training Dataset Count: " + str(train_data.count()))\nprint("Test Dataset Count: " + str(test_data.count()))
```

Training Dataset Count: 146904
Test Dataset Count: 62769

```
train_data.show()
```

trip_duration	distance_traveled	num_of_passengers	fare	tip	miscellaneous_fees	total_fare	surge_applied	features_scaled
0.0	0.02	1.0	150.0	0	2.25	152.25	0	(5,[1,2],[0.02,1.0
0.0	0.02	1.0	187.5	0	2.25	189.75	0	(5,[1,2],[0.02,1.0
0.0	0.06	2.0	225.0	45	2.6999999999999886	272.7	0	[0.0,0.06,2.0,45..
0.0	0.06	2.0	352.5	71	2.1999999999999886	425.7	0	[0.0,0.06,2.0,71..
0.0	0.1	1.0	60.0	8	1.75	69.75	0	[0.0,0.1,1.0,8.0,..
0.0	0.1	1.0	75.0	15	2.700000000000003	92.7	0	[0.0,0.1,1.0,15.0.
0.0	0.11	1.0	18.75	15	6.0	39.75	0	[0.0,0.11,1.0,15..
0.0	0.11	1.0	18.75	38	9.25	66.0	0	[0.0,0.11,1.0,38..
0.0	0.13	1.0	262.5	0	2.25	264.75	0	(5,[1,2],[0.13,1.0
0.0	0.16	1.0	165.0	1	1.8499999999999943	167.85	0	[0.0,0.16,1.0,1.0.
0.0	0.16	2.0	600.0	0	2.25	602.25	0	(5,[1,2],[0.16,2.0
0.0	0.18	1.0	18.75	38	9.25	66.0	0	[0.0,0.18,1.0,38..
0.0	0.18	1.0	187.5	38	2.1999999999999886	227.7	0	[0.0,0.18,1.0,38..
0.0	0.18	1.0	487.5	8	1.75	497.25	0	[0.0,0.18,1.0,8.0.
0.0	0.21	2.0	187.5	38	2.1999999999999886	227.7	0	[0.0,0.21,2.0,38..
0.0	0.23	1.0	487.5	98	2.2000000000000455	587.7	0	[0.0,0.23,1.0,98..
0.0	0.24	1.0	75.0	0	2.25	77.25	0	(5,[1,2],[0.24,1.0
0.0	0.4	1.0	60.0	16	1.8500000000000085	77.85000000000001	0	[0.0,0.4,1.0,16.0.
0.0	7.19	1.0	138.75	0	9.75	148.5	0	(5,[1,2],[7.19,1.0
1.0	0.02	1.0	18.75	375	6.0	399.75	0	[1.0,0.02,1.0,375.

only showing top 20 rows

▼ Training the Model

```
from pyspark.ml.regression import LinearRegression,DecisionTreeRegressor
from pyspark.ml.evaluation import RegressionEvaluator

# Initialize `lr`
lr = LinearRegression(labelCol="total_fare", maxIter=10, regParam=0.3, elasticNetParam=0.8)

# Fit the data to the model
linearModel = lr.fit(train_data)

# Generate predictions
predicted = linearModel.transform(test_data)
predicted.show()
```

trip_duration	distance_traveled	num_of_passengers	fare	tip	miscellaneous_fees	total_fare	surge_applied	features
0.0	0.02	1.0	18.75	0	13.5	32.25	0	(5,[1,2],[0.02,1.0]) (5,[1
0.0	0.02	1.0	135.0	0	2.25	137.25	0	(5,[1,2],[0.02,1.0]) (5,[1
0.0	0.05	1.0	45.0	0	2.25	47.25	0	(5,[1,2],[0.05,1.0]) (5,[1
0.0	0.43	1.0	300.0	60	2.6999999999999886	362.7	0	[0.0,0.43,1.0,60.... [0.0,
1.0	0.02	1.0	75.0	19	2.5999999999999943	96.6	0	[1.0,0.02,1.0,19.... [2.09
1.0	0.03	1.0	18.75	98	5.5	122.25	0	[1.0,0.03,1.0,98.... [2.09
1.0	0.03	1.0	112.5	29	1.9000000000000057	143.4	0	[1.0,0.03,1.0,29.... [2.09
1.0	0.05	1.0	150.0	30	2.6999999999999886	182.7	0	[1.0,0.05,1.0,30.... [2.09
1.0	0.05	2.0	75.0	15	2.7000000000000003	92.7	0	[1.0,0.05,2.0,15.... [2.09
1.0	0.11	1.0	18.75	0	6.0	24.75	0	[1.0,0.11,1.0,0.0.... [2.09
1.0	0.11	1.0	90.0	18	2.7000000000000003	110.7	0	[1.0,0.11,1.0,18.... [2.09
1.0	0.37	1.0	75.0	15	2.7000000000000003	92.7	0	[1.0,0.37,1.0,15.... [2.09
2.0	0.02	1.0	60.0	19	1.9249999999999972	80.925	0	[2.0,0.02,1.0,19.... [4.18
2.0	0.02	1.0	75.0	15	2.7000000000000003	92.7	0	[2.0,0.02,1.0,15.... [4.18
2.0	0.02	1.0	300.0	0	2.25	302.25	0	[2.0,0.02,1.0,0.0.... [4.18
2.0	0.03	1.0	18.75	0	6.0	24.75	0	[2.0,0.03,1.0,0.0.... [4.18
2.0	0.03	1.0	90.0	4	2.0	96.0	0	[2.0,0.03,1.0,4.0.... [4.18
2.0	0.03	1.0	150.0	0	2.25	152.25	0	[2.0,0.03,1.0,0.0.... [4.18
2.0	0.05	1.0	75.0	15	2.7000000000000003	92.7	0	[2.0,0.05,1.0,15.... [4.18
2.0	0.05	1.0	75.0	15	2.7000000000000003	92.7	0	[2.0,0.05,1.0,15.... [4.18

only showing top 20 rows

Getting the outcome

```
# Coefficients for the model
linearModel.coefficients

DenseVector([0.0026, 0.0152, 0.8405, 2.3339, 32.2252])

# Intercept for the model
linearModel.intercept

84.20828665780299

# Get the RMSE
linearModel.summary.rootMeanSquaredError

83.66711413638198

# Get the R2
linearModel.summary.r2

0.2939909414691463

dtr = DecisionTreeRegressor(labelCol="total_fare", featuresCol='features')

dtrModel = dtr.fit(train_data)

predicted = dtrModel.transform(test_data)
predicted.show()
```

trip_duration	distance_traveled	num_of_passengers	fare	tip	miscellaneous_fees	total_fare	surge_applied	features
0.0	0.02	1.0	18.75	0	13.5	32.25	0	(5,[1,2],[0.02,1.0]) (5,[1
0.0	0.02	1.0	135.0	0	2.25	137.25	0	(5,[1,2],[0.02,1.0]) (5,[1
0.0	0.05	1.0	45.0	0	2.25	47.25	0	(5,[1,2],[0.05,1.0]) (5,[1

8/15/23, 12:02 AMAbhishekLal_C23039 Project -II.ipynb - Colaboratory

	0.0	0.43	1.0	300.0	60	2.6999999999999886	362.7	0 [0.0,0.43,1.0,60.... [0.0,
	1.0	0.02	1.0	75.0	19	2.5999999999999943	96.6	0 [1.0,0.02,1.0,19.... [2.09
	1.0	0.03	1.0	18.75	98	5.5	122.25	0 [1.0,0.03,1.0,98.... [2.09
	1.0	0.03	1.0	112.5	29	1.9000000000000057	143.4	0 [1.0,0.03,1.0,29.... [2.09
	1.0	0.05	1.0	150.0	30	2.6999999999999886	182.7	0 [1.0,0.05,1.0,30.... [2.09
	1.0	0.05	2.0	75.0	15	2.7000000000000003	92.7	0 [1.0,0.05,2.0,15.... [2.09
	1.0	0.11	1.0	18.75	0	6.0	24.75	0 [1.0,0.11,1.0,0.0... [2.09
	1.0	0.11	1.0	90.0	18	2.7000000000000003	110.7	0 [1.0,0.11,1.0,18.... [2.09
	1.0	0.37	1.0	75.0	15	2.7000000000000003	92.7	0 [1.0,0.37,1.0,15.... [2.09
	2.0	0.02	1.0	60.0	19	1.9249999999999972	80.925	0 [2.0,0.02,1.0,19.... [4.18
	2.0	0.02	1.0	75.0	15	2.7000000000000003	92.7	0 [2.0,0.02,1.0,15.... [4.18
	2.0	0.02	1.0	300.0	0	2.25	302.25	0 [2.0,0.02,1.0,0.0... [4.18
	2.0	0.03	1.0	18.75	0	6.0	24.75	0 [2.0,0.03,1.0,0.0... [4.18
	2.0	0.03	1.0	90.0	4	2.0	96.0	0 [2.0,0.03,1.0,4.0... [4.18
	2.0	0.03	1.0	150.0	0	2.25	152.25	0 [2.0,0.03,1.0,0.0... [4.18
	2.0	0.05	1.0	75.0	15	2.7000000000000003	92.7	0 [2.0,0.05,1.0,15.... [4.18
	2.0	0.05	1.0	75.0	15	2.7000000000000003	92.7	0 [2.0,0.05,1.0,15.... [4.18
+-----+-----+-----+-----+-----+-----+-----+-----+-----+								
only showing top 20 rows								

```
evaluate_r2 = RegressionEvaluator(predictionCol="prediction",labelCol="total_fare",metricName="r2")

evaluate_r2.evaluate(predicted)

0.7179675648501878

evaluate = RegressionEvaluator(labelCol="total_fare",metricName="rmse")

rmse = evaluate.evaluate(predicted)
print("RMSE: ",rmse)

RMSE:  51.488353691131145
```